

Generative AI and LLM Foundations Executive Brief

Decision guide for responsible institutional adoption

Generative language models can accelerate drafting, summarization, extraction and knowledge work, but fluent output is not factual authority. Institutions should adopt them through bounded use cases, authorized evidence, measurable evaluation and accountable approval.

What leaders need to understand

Question	Executive answer
What does the model do?	Generates probable continuations from learned patterns and retained context.
What can it support?	Drafting, summarization, extraction, classification, translation and code assistance.
What can go wrong?	Unsupported claims, altered numbers, disclosure, bias, prompt injection and excessive authority.
What establishes trust?	Use-case evidence, source traceability, controls, monitoring and accountable human decisions.

Immediate decision rules

1. Start with a bounded, reversible and measurable workflow.
2. Keep authoritative data and deterministic calculations outside the language model.
3. Require claim-level verification for material facts.
4. Prohibit unapproved confidential data and uncontrolled tool access.
5. Measure accepted-output quality, incidents, reviewer effort, latency and total cost.
6. Assign an accountable owner with authority to stop or roll back deployment.

A practical control model

Stage	Required control	Evidence retained
Request	Purpose and user authorization	Task, user and intended audience
Context	Approved source and retrieval policy	Source IDs, versions and passages
Generation	Controlled configuration	Model, prompt, tools and decoding settings
Verification	Claim and risk checks	Exceptions, corrections and reviewer evidence
Decision	Human authorization	Approval, rejection or escalation
Operation	Monitoring and incident response	Performance, incidents and change record

Priority executive questions

1. Which decisions may the system influence, and which are excluded?
2. What source is authoritative when model text conflicts with institutional records?
3. How will affected people obtain correction or recourse?
4. Which changes trigger revalidation?
5. Who owns residual risk and incident response?

Learning and capability pathway

Book 4 is designed as a cumulative pathway. Across 30 lessons, ambitious learners will progress from tokens and next-token prediction through retrieval, tools, agents, fine-tuning, evaluation, safety and production architecture. The capstone will integrate these foundations into a small educational GPT-style system.

The promise is intentionally precise: learners can build and explain a compact decoder-only Transformer, train it on a modest corpus, generate text and evaluate its limitations. This is different from training a frontier model, which requires vastly larger datasets, compute, engineering teams and operational controls.

Executive takeaway

Adopt generative AI where evidence shows net value, preserve authoritative systems of record, and never confuse persuasive language with accountable institutional judgment.

SRAI Principle Number 96

A language model generates statistically plausible continuations; factual authority must come from evidence, validation and accountable human use.